

We will be importing the following packages to perform the analysis.

```
In [1]: 1 import numpy as np
        2 import pandas as pd
        3 import itertools
        4 from sklearn.model_selection import train_test_split
        5 from sklearn.feature_extraction.text import TfidfVectorizer
        6 from sklearn.linear_model import PassiveAggressiveClassifier
        7 from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
        8 import matplotlib.pyplot as plt
```

## Steps

1. Reading and Characterizing the Data (ACQUIRE)
2. Exploration (PREPARE)
3. Cleaning and Filtering the Data for our requirements (PREPARE)
4. Analysis (ANALYZE)
5. Results (REPORT)

## I. Characterizing the data (ACQUIRE)

Code below is to read the data file into a pandas dataframe.

```
In [2]: 1 df=pd.read_csv('news.csv')
```

Before we start exploring the data, we want a clearer picture of the data. Hence we ask a few questions.

### Q1. How many news records do we have?

Code below is to see the number of rows and columns in the data.

```
In [3]: 1 df.shape
```

```
Out[3]: (6335, 4)
```

We can see that the dataframe has 4 columns and 6335 records .

## Q2. What are the column names?

Code below is to see the column names of the dataframe

```
In [4]: 1 df.columns
```

```
Out[4]: Index(['Unnamed: 0', 'title', 'text', 'label'], dtype='object')
```

We can see that only 3 out of 4 columns have names. They are title , text , label .

## Q3. What does the data look like?

Code below is to see a sample of the data.

```
In [5]: 1 #df.head()  
2 df.tail()
```

```
Out[5]:
```

	Unnamed: 0		title		text	label
6330	4490		State Department says it can't find emails fro...		The State Department told the Republican Natio...	REAL
6331	8062		The 'P' in PBS Should Stand for 'Plutocratic' ...		The 'P' in PBS Should Stand for 'Plutocratic' ...	FAKE
6332	8622		Anti-Trump Protesters Are Tools of the Oligarc...		Anti-Trump Protesters Are Tools of the Oligar...	FAKE
6333	4021		In Ethiopia, Obama seeks progress on peace, se...		ADDIS ABABA, Ethiopia —President Obama convene...	REAL
6334	4330		Jeb Bush Is Suddenly Attacking Trump. Here's W...		Jeb Bush Is Suddenly Attacking Trump. Here's W...	REAL

Now we have understood the basic characteristics of the data. Let's proceed with the next step.

## II. Exploration (PREPARE)

Code below is to create a new dataframe that contains only the labels(REAL/FAKE).

```
In [6]: 1 labels = df.label
```

Code below is to see a sample of the newly created labels dataframe.

```
In [7]: 1 labels.head()
```

```
Out[7]: 0    FAKE
        1    FAKE
        2    REAL
        3    FAKE
        4    REAL
        Name: label, dtype: object
```

Code below is to split the dataset into training and testing sets using the train\_test\_split function.

```
In [8]: 1 x_train, x_test, y_train, y_test = train_test_split(df['text'], labels, test_size=0.3, random_state=7)
```

Code below is to see the training datasets created using the train\_test\_split function.

```
In [9]: 1 x_train
        2 #y_train
```

```
Out[9]: 4274    Home / Be The Change / Government Corruption /...
        4310    At least a half-dozen attendees shoved and tac...
        2050    As soon as Rep. Kevin McCarthy (R-Calif.) shoc...
        4410    NOT ON THE SHORT LIST\n\nFormer Vice President...
        3106    Eric Liu is the founder of Citizen University ...
              ...
        5699
        2550    It's not that Americans won't elect wealthy pr...
        537     Anyone writing sentences like 'nevertheless fu...
        1220    More Catholics are in Congress than ever befor...
        4271    It was hosted by CNN, and the presentation was...
        Name: text, Length: 4434, dtype: object
```

Code below is to see the testing datasets created using the `train_test_split` function.

```
In [10]: 1 #x_test  
        2 y_test
```

```
Out[10]: 3534    REAL  
        6265    FAKE  
        3123    REAL  
        3940    REAL  
        2856    REAL  
        ...  
        118     FAKE  
        3258    REAL  
        4521    FAKE  
        5926    FAKE  
        89      REAL  
Name: label, Length: 1901, dtype: object
```

### III. Cleaning and Filtering the Data for our requirements (PREPARE)

Stop words are the most common words in a language that are considered to be useless and is mostly filtered out before processing the natural language data. A simple way to filter out stopwords is to just use the `corpus` from `nltk` that we can download easily and remove the stopwords from our text using a `loop`. Let's use a different approach this time.

#### Term Frequency – Inverse Document Frequency

An alternative is to calculate word frequencies, is called TF-IDF.

This is an acronym that stands for “Term Frequency – Inverse Document Frequency” which are the components of the resulting scores assigned to each word.

- Term Frequency : This summarizes how often a given word appears within a document.
- Inverse Document Frequency : This downscales words that appear a lot across documents.

To simply state, TF-IDF are word frequency scores that try to highlight words that are interesting, rather than merely highlighting the most frequent words. We can omit words which cross a certain `threshold(frequency)`.

- A `TfidfVectorizer` turns a collection of raw documents into a matrix of TF-IDF features

- `max_df` is the parameter for threshold that we give as an input for the function.
- `max_df` can be set to a value in the range `[0.7, 1.0]` to automatically detect and filter stop words based on intra corpus document frequency of terms

Example: Stopwords like `the` is very commonly used and will most likely have document frequency higher than `0.7` in a given sample news text.

Let's initialize a `TfidfVectorizer` with maximum document frequency of `0.7` (terms with a higher document frequency will be discarded).

```
In [11]: 1 vectorizer = TfidfVectorizer(max_df=0.7)
```

The `fit` method, when applied to the training dataset, learns the model parameters. We should then apply the `transform` method on the training dataset to get the transformed (scaled) training dataset. Instead of performing them individually, we perform both of these steps in one step by applying `fit_transform` on the training dataset

Code below is to `fit` and `transform` the vectorizer on the training set.

```
In [12]: 1 tfidf_train = vectorizer.fit_transform(x_train)
```

But for testing set, Machine Learning applies prediction based on what was learned during the fitting of the training set and so it doesn't need to learn the models parameters, it directly performs the transformation.

Code below is to `transform` the vectorizer on the test set.

```
In [13]: 1 tfidf_test = vectorizer.transform(x_test)
```

Code below to is to see the transformed training and test datasets.

In [14]:

```
1 #print(tfidf_train)
2 print(tfidf_test)

(0, 57224)    0.05519528105096844
(0, 57178)    0.031192704431556396
(0, 56860)    0.023712653791672828
(0, 56793)    0.04250547812724162
(0, 56783)    0.07367402497455956
(0, 56305)    0.0977655117405021
(0, 56197)    0.029044530183008413
(0, 56172)    0.024490916832987555
(0, 56150)    0.023453885650831386
(0, 56080)    0.05048930873591531
(0, 55877)    0.1127383261667559
(0, 55672)    0.10262396941875963
(0, 55601)    0.07458532488029349
(0, 55239)    0.06591601888097635
(0, 54939)    0.06283196078414083
(0, 54424)    0.030434000427700194
(0, 54292)    0.02428650983920786
(0, 54226)    0.04242613384089784
(0, 53748)    0.11573181749112645
(0, 51943)    0.19609860066137222
(0, 51623)    0.03358531291902588
(0, 51592)    0.038177147673052726
(0, 51497)    0.03312197768596556
(0, 51470)    0.02148967190615234
(0, 51426)    0.02362395217837374
:
(1900, 4928)  0.019341505823414613
(1900, 3804)  0.014078619628178721
(1900, 3767)  0.02002705524036257
(1900, 3495)  0.015979793733507684
(1900, 3433)  0.02987833791332498
(1900, 3290)  0.01008872712538997
(1900, 3267)  0.045897743854312854
(1900, 3259)  0.018458331365985684
(1900, 3213)  0.06690993887860666
(1900, 3168)  0.009454303787877413
(1900, 2887)  0.027246814620999194
(1900, 2867)  0.05397074111861263
```

```

(1900, 2838) 0.02148166745946232
(1900, 2815) 0.034435373095004655
(1900, 2778) 0.0323437430969293
(1900, 2756) 0.024703404395062703
(1900, 2737) 0.01624512944973044
(1900, 2735) 0.037969113344949915
(1900, 2721) 0.01048730126970561
(1900, 2317) 0.01735203645639521
(1900, 1970) 0.07328349111234908
(1900, 1020) 0.04744390200517566
(1900, 631) 0.021224832383165613
(1900, 273) 0.019881514371597356
(1900, 1) 0.017138061540192425

```

Code below is to see the words which had a document frequency less than 0.7 .

In [15]: 1 print(vectorizer.vocabulary\_)

```

40214, 'foreign': 20618, 'until': 54226, 'proven': 40816, 'otherwise': 36924, 'bad': 5454, 'none': 35649, 'neat': 350
07, 'colorblind': 11125, 'ultimately': 53270, 'deepest': 14027, 'racial': 41541, 'start': 48864, 'lazy': 29906, 'mess
y': 32888, 'historical': 24541, 'truths': 52892, 'stories': 49272, 'story': 49281, 'endings': 17683, 'cheap': 9939,
'point': 39413, 'mlk': 33631, 'redemption': 42440, 'easy': 16887, 'honor': 24819, 'official': 36374, 'commissions': 1
1287, 'conversations': 12159, 'ourselves': 36960, 'washington': 55765, 'continues': 12057, 'force': 20577, 'heading':
23902, 'election': 17222, 'orc': 36737, 'poll': 39516, 'broad': 8087, 'lead': 29925, 'field': 19839, 'democratic': 14
353, 'challengers': 9752, 'nomination': 35615, 'contest': 12031, 'sizable': 47258, 'contenders': 12020, 'side': 4694
7, 'general': 21742, 'match': 32175, 'ups': 54352, 'candidates': 8909, 'gets': 21924, 'within': 56609, '10': 109, 'po
ints': 39422, 'hypothetical': 25393, 'matchups': 32180, 'rand': 41778, 'closest': 10752, '43': 1100, 'likely': 30474,
'54': 1266, 'walker': 55602, 'equally': 18053, 'carrying': 9218, '55': 1275, 'huckabee': 25115, '41': 1080, 'carson':
9221, '56': 1282, 'warren': 55740, 'decide': 13891, 'stands': 48811, 'benefit': 6429, 'gaining': 21436, 'holding': 24
674, '67': 1410, '16': 312, 'advantage': 2557, 'biden': 6708, 'backers': 5400, 'allocated': 3219, 'choice': 10196, 'n
otably': 35808, 'surges': 50065, '74': 1485, 'broadly': 8100, 'believe': 6339, 'chances': 9778, 'hold': 24670, 'stron
gest': 49466, '68': 1415, 'better': 6609, 'shot': 46772, 'ticket': 51714, 'favorability': 19470, 'rating': 41917, 're
cently': 42244, 'prospects': 40741, 'appear': 4059, 'unchanged': 53423, 'compared': 11357, 'polls': 39529, 'conducte
d': 11624, 'broke': 8123, 'personal': 38528, 'email': 17389, 'address': 2391, 'based': 5907, 'server': 46179, 'servin
g': 46192, 'leads': 29935, 'pack': 37405, 'follows': 20502, '13': 226, 'nearly': 35004, 'matches': 32177, '12': 199,
'holds': 24680, 'backing': 5410, 'dipped': 15261, 'significantly': 47027, 'february': 19563, 'generally': 21750, 'ste
ady': 48974, 'single': 47163, 'digits': 15154, 'jersey': 27881, 'gov': 22489, 'christie': 10268, 'sen': 46034, 'marc

```

If we explore the words printed above, we can notice that stopwords have been removed automatically.

## IV. Analysis (ANALYZE)

### Passive Aggressive Classifier

This type of classifier is generally used for large-scale learning. It is one of the few online-learning algorithms. In online machine learning algorithms, the input data comes in sequential order and the machine learning model is updated step-by-step, as opposed to batch learning, where the entire training dataset is used at once.

- **Passive** : If the prediction is correct, keep the model and do not make any changes. i.e., the data in the example is not enough to cause any changes in the model.
- **Aggressive** : If the prediction is incorrect, make changes to the model. i.e., some change to the model may correct it.
- The parameter `max_iter` denotes the maximum number of iterations the model makes over the training data.

This is very useful in situations where there is a huge amount of data and it is computationally infeasible to train the entire dataset because of the sheer size of the data. We can simply say that an online-learning algorithm will get a training example, update the classifier, and then throw away the example.

A very good example of this would be to detect fake news on a social media website like Twitter, where new data is being added every second. To dynamically read data from Twitter continuously, the data would be huge, and using an online-learning algorithm would be the ideal choice.

First, we initialize a `PassiveAggressiveClassifier`. Then we fit this on `tfidf_train` and `y_train`. Post fitting, we predict on the test set.

Code below to initialize a Passive Aggressive Classifier.

```
In [16]: 1 pac = PassiveAggressiveClassifier(max_iter=50)
```

Code below to fit the classifier using the training datasets.

```
In [17]: 1 pac.fit(tfidf_train,y_train)
```

```
Out[17]: PassiveAggressiveClassifier(C=1.0, average=False, class_weight=None,
    early_stopping=False, fit_intercept=True,
    loss='hinge', max_iter=50, n_iter_no_change=5,
    n_jobs=None, random_state=None, shuffle=True,
    tol=0.001, validation_fraction=0.1, verbose=0,
    warm_start=False)
```



We can see above that we have created a `Passive Aggressive Classifier` .

Code below to `predict` on the test set from the `TfidfVectorizer`

```
In [18]: 1 y_pred = pac.predict(tfidf_test)
```

Calculate the accuracy with `accuracy_score()` from `sklearn.metrics` .

```
In [19]: 1 score=accuracy_score(y_test,y_pred)
```

## V. Results (REPORT)

Code below to assign the accuracy value to a variable for convenience.

```
In [20]: 1 d = round(score*100,2)
```

Code below to print the accuracy of the model.

```
In [21]: 1 print("Accuracy: ",d,"%")
```

Accuracy: 92.9 %

Confusion matrix is used to display the exact predictions of a classifier model.

Code below to build a confusion matrix

```
In [22]: 1 cmatrix = confusion_matrix(y_test,y_pred, labels=['FAKE','REAL'])
```

Code below to print the confusion matrix.

```
In [23]: 1 print(cmatrix)
```

```
[[904  70]
 [ 65 862]]
```

To better understand the confusion matrix, we will access the matrix and display the data in a easily understandable manner.

Code below to do the above mentioned actions.

```
In [24]: 1 w = cmatrix[0][0]
        2 x = cmatrix[0][1]
        3 y = cmatrix[1][1]
        4 z = cmatrix[1][0]
        5 print("True Positives =",w)
        6 print("False Positives =",x)
        7 print("True Negatives =",y)
        8 print("False Negatives =",z)
```

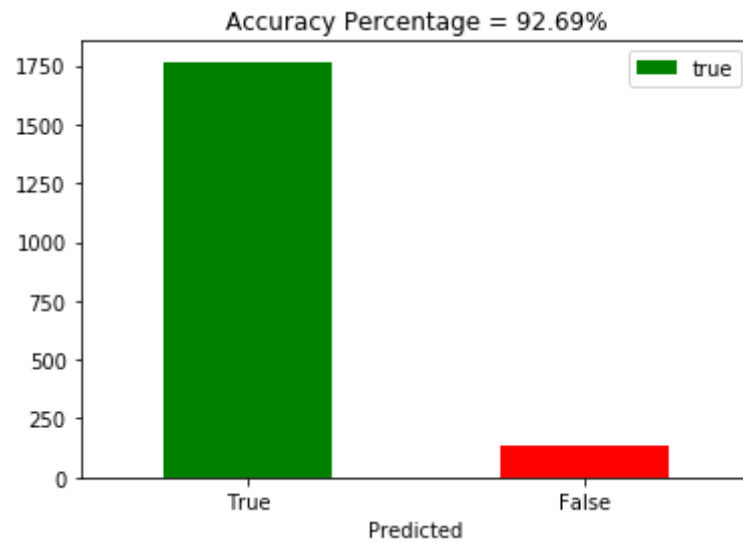
```
True Positives = 904
False Positives = 70
True Negatives = 862
False Negatives = 65
```

In order to further understand our findings, we can vizualise it.

Code below is to vizualise our findings.

In [25]:

```
1 a = w + y
2 b = x + z
3 df = pd.DataFrame({'Predicted': ['True', 'False'], 'true': [a, b]})
4
5 ax = df.plot.bar(x='Predicted', y='true', rot=0, title="Accuracy Percentage = 92.69%", color=['green', 'red'])
```



Code below to print the conclusion.

In [26]: 1 `print("So with this model, we have ",w+y," correctly predicted labels, and ",x+z," wrongly predicted labels.")`

So with this model, we have 1766 correctly predicted labels, and 135 wrongly predicted labels.

In [ ]: 1